


# tutorial5

Likun Cui

04 September, 2018

## Partition the data

```
# create positive class sample with 2 descriptive features set.seed(3)
f1 <- rnorm(100, mean=6, sd = 1.2)
set.seed(4)
f2 <- rnorm(100, mean=6, sd = 1.2)
P.data <- cbind(f1, f2)
# create positive class sample with 2 descriptive features set.seed(7)
f1 <- rnorm(300, mean=4, sd = 1.2)
set.seed(8)
f2 <- rnorm(300, mean=4, sd = 1.2)
N.data <- cbind(f1, f2)
# combine all samples
data.mat <- data.frame(rbind(P.data, N.data), Class=rep(c(1, 0), time=c(nrow(P.data), nrow(N.data))))
plot(subset(data.mat, Class==1)[,-3], col="red", pch=16, ylim=c(0, 9), xlim=c(0, 9), xlab="Feature 1", ylab="Feature 2")
points(subset(data.mat, Class==0)[,-3], col="blue", pch=16)
```



```
dim(N.data)
```

```
## [1] 300  2
```

```
sub<-sample(1:nrow(N.data),round(nrow(N.data)*0.8))  
length(sub)
```

```
## [1] 240
```

```
data_train<-N.data[sub,]  
data_test<-N.data[-sub,]  
dim(data_train)
```

```
## [1] 240  2
```


```
dim(data_test)
```

```
## [1] 60  2
```

```
#data_train
#print(data)
#data.mat
#data_train.mat
```


# Train a Logistic Regression

```
# train a logistic regression model
logit.model <- glm(Class~., family=binomial(link='logit'), data=data.mat)
# plot fitted values from logistic regression model
plot(logit.model$fitted.values)
```



## Train an LDA

```
library(MASS)
# train an LDA model
lda.model <- lda(Class~., data=data.mat)
lda.fitted <- predict(lda.model, data.mat)$posterior[, "1"]
# plot fitted values from LDA model
plot(lda.fitted)
```



```
# use fitted value to classify samples
lda.decision <- ifelse(lda.fitted > 0.5, 1, 0)
# calculate classification accuracy (in percentage %)
sum(lda.decision == data.mat$Class) / nrow(data.mat) * 100
```

```
## [1] 92.5
```

## Train an KNN

```
library(class)
# a knn with k=1
knn.model1 <- knn(train=data.mat[, -3], test=data.mat[, -3], cl=data.mat[, 3], k=1)
# a knn with k=5
knn.model2 <- knn(train=data.mat[, -3], test=data.mat[, -3], cl=data.mat[, 3], k=5)
# a knn with k=50
knn.model3 <- knn(train=data.mat[, -3], test=data.mat[, -3], cl=data.mat[, 3], k=50)

# calculate classification accuracy
sum(knn.model1 == data.mat$Class) / nrow(data.mat) * 100
```

```
## [1] 100
```


```
sum(knn.model2 == data.mat$Class) / nrow(data.mat) * 100
```

```
## [1] 93.5
```

```
sum(knn.model3 == data.mat$Class) / nrow(data.mat) * 100
```

```
## [1] 92.5
```

```
# apply knn and enable calculation of prediction probability  
knn.model <- knn(train=data.mat[, -3], test=data.mat[, -3], cl=data.mat[, 3], k=5, prob=TRUE)  
# extract prediction probability from the prediction model  
knn.prob <- attr(knn.model, "prob")  
isNegativeSample <- data.mat[, 3] != 1  
knn.prob[isNegativeSample] <- 1 - knn.prob[isNegativeSample]  
  
plot(knn.prob)
```

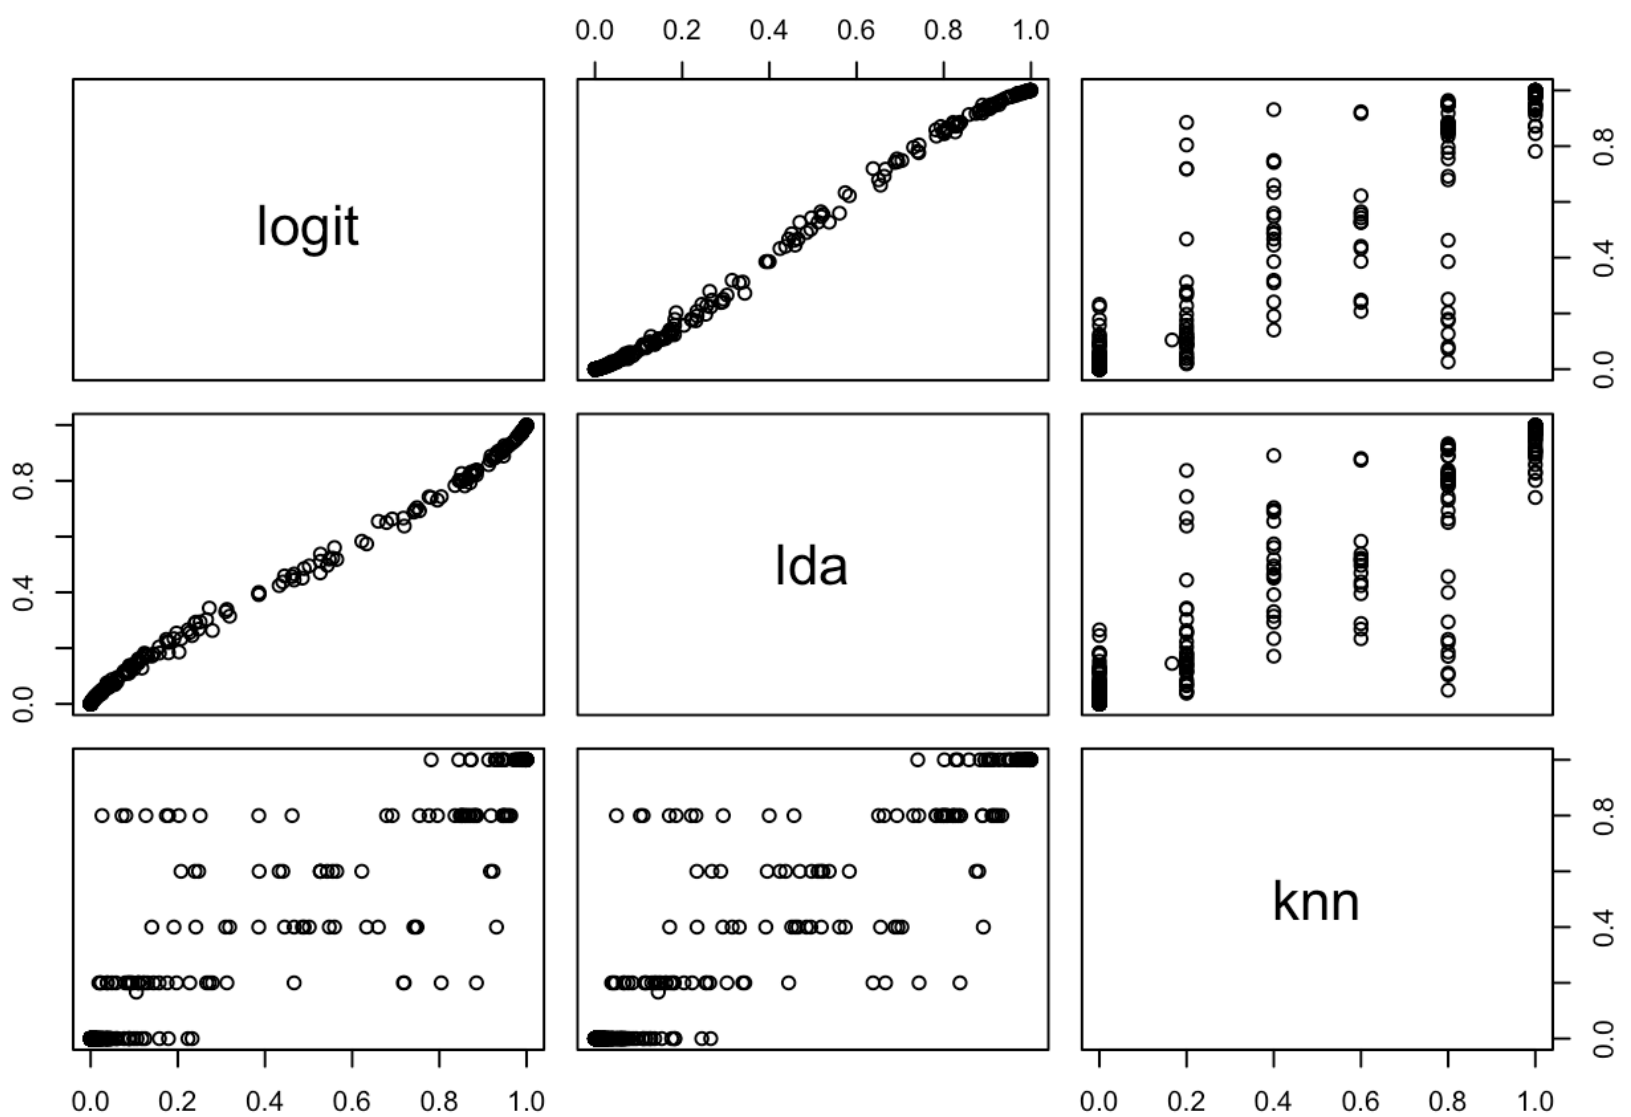


# Compare their performance

```
# combine classification results into a data frame
classifications <- data.frame(logit=logit.model$fitted.values, lda=lda.fitted, knn=
=knn.prob)
# calculate correlation
cor(classifications)
```

```
##           logit      lda      knn
## logit 1.0000000 0.9986088 0.9284626
## lda   0.9986088 1.0000000 0.9339678
## knn   0.9284626 0.9339678 1.0000000
```

```
# create pairwise scatter plot
pairs(classifications)
```



# identify optimal k value by minimising classification error on test set.

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
set.seed(1)
inTrain <- createDataPartition(data.mat$Class, p = .8)[[1]]
dataTrain <- data.mat[ inTrain, ]
dataTest  <- data.mat[-inTrain, ]
set.seed(1)
KNN1 <- train(Class~f1,
              data = dataTrain,
              method = "knn",
              trControl = trainControl(method = "repeatedcv",
                                       repeats = 5))
```

```
## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to
## do classification? If so, use a 2 level factor as your outcome column.
```

```
## Print diagnostic and summary information and statistics fo the model
KNN1
```

```
## k-Nearest Neighbors
##
## 320 samples
## 1 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 288, 288, 288, 288, 288, 288, ...
## Resampling results across tuning parameters:
##
##  k  RMSE          Rsquared    MAE
##  5  0.3351509    0.4249579    0.1982500
##  7  0.3242995    0.4516295    0.1987500
##  9  0.3272536    0.4417977    0.2040833
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 7.
```

```
print("Therefore, k=9 is the best choice.")
```

```
## [1] "Therefore, k=9 is the best choice."
```

**Now we used test set to select optimal k, is it still valid to use this test set to evaluate the performance of our optimised kNN classifier?**

# Why or why not?

```
print("No. Since the chosen k would be perfectly predicting the test set, a new data set should be developed to evaluate the performance.")
```

```
## [1] "No. Since the chosen k would be perfectly predicting the test set, a new data set should be developed to evaluate the performance."
```